

Casey Bray

Rapid City, South Dakota · 605-787-8604 · cbrayanalytics@gmail.com
[linkedin.com/in/caseylbray](https://www.linkedin.com/in/caseylbray) · github.com/cbrayanalytics · cbrayanalytics.github.io/portfolio

SUMMARY

Data professional bridging production infrastructure and analytics. 4+ years operating Linux and distributed-storage systems — Cassandra, Redis, AWS S3 — at petabyte scale, paired with an M.S. in Data Analytics and hands-on Python/R modeling. I keep data systems reliable *and* turn the data they hold into decisions.

SKILLS

Data & Analytics: Python, R, SAS, SQL · pandas, scikit-learn · statistical modeling, machine learning · Tableau

Infrastructure & Engineering: Linux · Spark, Hadoop, Hive · Cassandra (CQL), Redis · AWS S3 · Ansible, Terraform · shell scripting · Grafana · log analysis

EXPERIENCE

Cloudian — Critical Case Engineer

Feb 2025 – Present

- Partner with engineering to resolve **~5 critical production escalations per week** that exceed standard support, driving root-cause fixes for issues regular support cannot.
- Diagnose root cause by analyzing time-series **Grafana** metrics and parsing system logs with shell tooling (grep/awk/sed) — correlating request latency, throughput, and error patterns into customer-facing guidance.
- Query and maintain **Cassandra** clusters via **CQL**: validate data consistency, tune compactions and snapshot cleanup, and remediate records to control database size and system performance.
- Build custom shell scripts to automate workflows and remediate critical bugs.

Cloudian — Principal L2 Technical Support Engineer

Aug 2023 – Feb 2025

- Resolved the most complex escalations requiring multi-phase remediation — manipulating metadata and object data, modifying code, and writing custom shell scripts for petabyte-scale S3 deployments.
- Prevented loss of **1+ petabyte** of medical imaging data by diagnosing and correcting an inefficient storage configuration at a Midwest storage provider.
- Surfaced **10+ bugs and missing features** later shipped in HyperStore S3 updates; documented each in JIRA with reproduction data for engineering.

Cloudian — Managing Technical Operations Engineer

Oct 2022 – Aug 2023

- Led and mentored a global team of **8 support engineers** through technical training and live customer escalations.
- Built and maintained AWS cluster automation using **Python, Ansible, Terraform, and shell scripting**.

Cloudian — Technical Operations Engineer

Dec 2021 – Oct 2022

- Supported distributed storage and backup systems — AWS S3, Cassandra, Redis, Linux networking and hardware — via Salesforce, phone, and Zoom.
- Authored **15+ Confluence guides** on Linux administration, shell scripting, Cassandra, and Redis, raising team resolution efficiency.

WL Plastics — CDL Class A Resin Transportation Specialist

Mar 2020 – Dec 2021

- Cut offload times by **~15 min/trip** with a one-way valve filter eliminating contamination and repetitive setup; reduced forklift-collision insurance claims by **~50%** via steel safety pylons.

EDUCATION

Dakota State University — M.S., Data Analytics (Dec 2022)

Coursework: Modern Applied Statistics, Statistical Programming, Predictive Analytics, Data Mining, Data Warehousing. Python, R, SAS, Spark, Hadoop, Hive, SQL.

National American University — B.S., Business Management (Dec 2015)

SELECTED PROJECTS — FULL PORTFOLIO: [CBRAYANALYTICS.GITHUB.IO/PORTFOLIO](https://cbrayanalytics.github.io/portfolio)

- **Molecular Toxicity Classification (R)** — RFE reduced 1,203 descriptors to 8 with no accuracy loss; Random Forest.
- **Auto Loan Default Prediction (Python)** — XGBoost on 40+ features, imbalanced target; 78% test accuracy.
- **Salary Analysis of Data-Related Jobs (Tableau)** — 7-view interactive dashboard on 27,741 visa wage filings.